

Baseline Model

August 30, 2018

1 Firts attempts & Base Model

```
In [1]: from pathlib import Path
import pandas as pd
import numpy as np
from datetime import datetime
import time
import matplotlib.pyplot as plt
%matplotlib inline
#%pylab inline
import itertools
import pickle
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.decomposition import PCA
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, f1_score, precision_score, recall_score
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import LabelBinarizer
from sklearn.preprocessing import OneHotEncoder
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from sklearn.feature_selection import RFECV, RFE

In [2]: path_project = Path.home() / Path('Google Drive/Felix')
path_data = path_project / Path("data")
path_dump = path_project / Path("dump")

In [3]: # loading cdv data
file = path_data / Path("felix.csv")
with Path.open(file, 'rb') as fp:
    cdv = pd.read_csv(fp, encoding='cp1252', low_memory=False)

In [4]: # loadind cdv data without format
file = path_data / Path("felix_ssfmt.csv")
with Path.open(file, 'rb') as fp:
    cdv_ssfmt = pd.read_csv(fp, encoding='cp1252', low_memory=False)
```

1.1 1) Feature engineering

```
In [5]: filename = path_dump / Path("dict_var_groups.sav")
        with open(filename, 'rb') as fp:
            dict_var_groups = pickle.load(fp)
```

```
scope_2015_var = dict_var_groups['scope_2015_var']
scope_2016_var = dict_var_groups['scope_2016_var']
scope_2017_var = dict_var_groups['scope_2017_var']
scope_2018_var = dict_var_groups['scope_2018_var']
scope_2015_2018_var = dict_var_groups['scope_2015_2018_var']
scope_2016_2018_var = dict_var_groups['scope_2016_2018_var']
scope_2017_2018_var = dict_var_groups['scope_2017_2018_var']
pred_var = dict_var_groups['pred_var']
com_var = dict_var_groups['com_var']
tech_var = dict_var_groups['tech_var']
text_var = dict_var_groups['text_var']
bizz_var = dict_var_groups['bizz_var']
cat_var = dict_var_groups['cat_var']
cat_max9_var = dict_var_groups['cat_max9_var']
cat_min10_var = dict_var_groups['cat_min10_var']
quant_var = dict_var_groups['quant_var']
```

```
In [6]: print(f"out of the {cdv.shape[1]} variable :")
        print(f"{len(scope_2015_2018_var)} variables como, to all years ")
```

out of the 354 variable :
268 variables como, to all years

```
In [7]: exclusion = com_var | tech_var | bizz_var | text_var
        print(f"Out of the {len(scope_2015_2018_var)} variables comon to all years ")
        print(f"{len(scope_2015_2018_var & exclusion)} are excluded ")
        scope_2015_2018_var_kept = scope_2015_2018_var - exclusion
        print(f"{len(scope_2015_2018_var_kept)} are kept ")
```

Out of the 268 variables comon to all years
22 are excluded
246 are kept

```
In [8]: print(f"out of the {len(scope_2015_2018_var_kept)} common variable :")
        print(f"{len(cat_var & scope_2015_2018_var_kept)} variables are categorial ")
        print(f"{len(quant_var & scope_2015_2018_var_kept)} variables are quantitative ")
```

out of the 246 common variable :
171 variables are categorial
75 variables are quantitative

```
In [9]: print(f"out of the {len(cat_var & scope_2015_2018_var_kept)} variable categorical:")
        print(f"{len(cat_max9_var & scope_2015_2018_var_kept)} variables have maximum 9 modaliti")
        print(f"{len(cat_min10_var & scope_2015_2018_var_kept)} variables have more and are excl")
        cat_var_kept = cat_max9_var & scope_2015_2018_var_kept
```

out of the 171 variable categorical:
156 variables have maximum 9 modalities
15 variables have more and are excluded

```
In [10]: scope_quant_var = (quant_var & scope_2015_2018_var_kept)
         quant_null = np.sum(cdv_ssfmt.loc[:,scope_quant_var].isnull())
         quant_var_kept = set(quant_null[quant_null < 200].index)
         print(f"out of the {len(scope_quant_var)} quantitatives variables:")
         print(f"{len(quant_var_kept)} have less than 200 missing values and are kept")
```

out of the 75 quantitatives variables:
60 have less than 200 missing values and are kept

```
In [11]: scope = cat_var_kept | quant_var_kept
         df = cdv_ssfmt.loc[:,scope]
         df.loc[:,cat_var_kept - {"HEUREUX"}] = cdv.loc[:,cat_var_kept - {"HEUREUX"}]
```

```
In [12]: print(f"Number of variable kept {df.shape[1]}")
```

Number of variable kept 216

1.1.1 Encoding 'HEUREUX'

```
In [13]: # reducing problem to a 2 class classification problem
         df["HEUREUX_CLF"] = 0
         df.loc[df["HEUREUX"]==4, "HEUREUX_CLF"] = 1
         df.loc[df["HEUREUX"]==3, "HEUREUX_CLF"] = 1
         df.loc[df["HEUREUX"]==5, "HEUREUX_CLF"] = None
```

```
In [14]: # Modelisation as a regression problem
         df["HEUREUX_REG"] = df["HEUREUX"]
         df.loc[df["HEUREUX"]==5, "HEUREUX_REG"] = None
```

1.1.2 Encoding categorical variables

```
In [15]: p = df.shape[1]
         print(f"{p} columns")
         print(f"out of which {len(cat_var_kept)-1} are corresponding to categorical features")
```

218 columns
out of which 155 are corresponding to categorical features

```
In [16]: df = pd.get_dummies(df,
                             columns=cat_var_kept - {"HEUREUX"},
                             dummy_na = True,
                             drop_first=1)

In [17]: q = df.shape[1]
         print(f"{q} columns after encoding")
         print(f"{len(cat_var_kept)-1} variables where re-encoded in {len(cat_var_kept)-1+q-p}")

635 columns after encoding
155 variables where re-encoded in 572
```

```
In [18]: def get_related_features(variable):
         '''return all columns in global variable df
         starting by variable'''
         if isinstance(variable, str):
             scope = {variable}
         else:
             scope = set(variable)
         features = set()
         for element in scope:
             features = features | {c for c in df.columns if len(c) > len(element)
                                     and c[0:len(element)] == element
                                     and c[len(element)]=='_'}
         return features
```

1.2 2) Dataset construction and visualisation

```
In [19]: # subsetting dataframe
         features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
         pred = "HEUREUX_CLF"

         # treating remaining missing values
         df_tmp = df.loc[:,set(features) | {pred} ].dropna()

         X = df_tmp.loc[:,features]
         y = df_tmp[pred]

         X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                             test_size=0.2,
                                                             random_state=42)

         scaler = StandardScaler().fit(X_train)
         X_train = scaler.transform(X_train)
         X_test = scaler.transform(X_test)

         print(f"Number exemple: {y.shape[0]}\n- training set: \
               {y_train.shape[0]}\n- test set: {y_test.shape[0]}")
```

```

print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

Number exemple: 10596

- training set: 8476

- test set: 2120

Number of features: p=632

Number of class: 2

class 0 : 34.9%

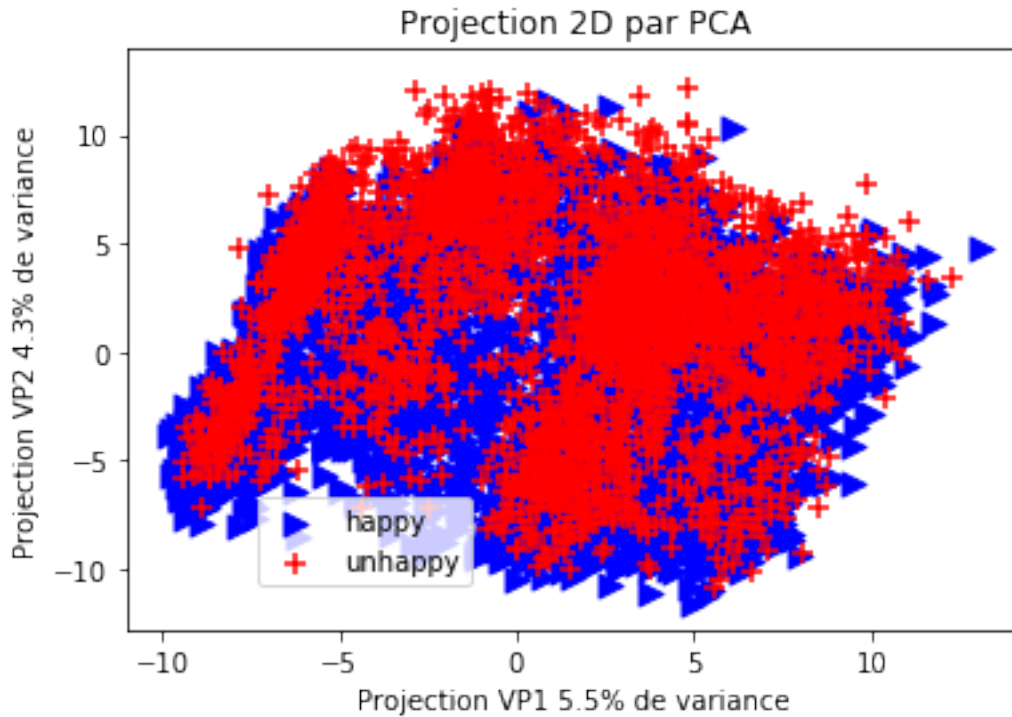
class 1 : 65.1%

In [20]: *# Reduction dim PCA*

```

pca = PCA(n_components=2)
pca.fit(X_train)
X_r = pca.transform(X_train)
happy = (y_train==1)
unhappy = (y_train==0)
plt.scatter(X_r[happy,0], X_r[happy,1],
            s=80, c='blue',marker=">", label="happy")
plt.scatter(X_r[unhappy,0], X_r[unhappy,1],
            s=80, c='red',marker='+', label="unhappy")
plt.ylabel(f'Projection VP2 \
{pca.explained_variance_ratio_[1]*100:0.1f}% de variance')
plt.xlabel(f'Projection VP1 \
{pca.explained_variance_ratio_[0]*100:0.1f}% de variance')
plt.title("Projection 2D par PCA")
plt.legend(bbox_to_anchor=(0.4, 0.25))
plt.show()

```



1.3 3) Baseline model

Logistic regression with 4 features to predict a 2 class variable

1.3.1 a) Training set and test set preparation

```
In [21]: # subsetting dataframe
features = {'ETATSAN', 'NOT_AMIS', 'DEPLOY', 'NOT_FAMI'}
pred = "HEUREUX_CLF"

# treating remaining missing values
df_tmp = df.loc[:,set(features) | {pred} ].dropna()

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
```

```

print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

Number exemple: 10915

- training set: 8732

- test set: 2183

Number of features: p=4

Number of class: 2

class 0 : 34.9%

class 1 : 65.1%

1.3.2 b) Hyperparameters tuning

```

In [22]: nb_value = 50 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
C_log = np.logspace(-2,2,nb_value)
cv = 6 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
std_scores_l1 = np.empty(nb_value)

np.random.seed(seed=42)

startTime = time.time()


for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1',
                             tol=0.01, random_state=42,
                             class_weight='balanced')
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))
    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))

plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:]-0.5*std_scores_l1[:],
              'r--', label=u'+/-0.5 écart type')

```

```
plt.semilogx(C_log,mean_score_l1[:]+0.5*std_scores_l1[:],'r--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C \npar validation croisée \n(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))
```



Détermination des paramètres optimaux en 17.5 s
Pénalisation l1, valeur optimale : C = 0.29

1.3.3 c) Model training and evaluation

```
In [23]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
```



```

        penalty='l1',
        tol=0.01,
        random_state=42,
        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Model score

- Accuracy : 67.1 %
- Precision : 76.2 % (Happy # positive class)
- Recall : 71.9 %
- F1 score : 74.0 %

1.4 4) Model selection (2 class)

1.4.1 a) Finding optimal set of features by recursive elimination using "Lasso" (RFECV)

```

In [24]: # subsetting dataframe
features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
pred = "HEUREUX_CLF"

# treating remaining missing values
print(f"{df.shape[0]} exemples before dropping na")
df_tmp = df.loc[:,set(features) | {pred} ].dropna()

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42
                                                    )

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")

```

```

print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

11131 exemples before dropping na

Number exemple: 10596

- training set: 8476

- test set: 2120

Number of features: p=632

Number of class: 2

class 0 : 34.9%

class 1 : 65.1%

In [25]: startTime = time.time()

```

scoring='accuracy'
step = 0.05

```

```

clf = LogisticRegression(C=1,
                        penalty='l1',
                        class_weight='balanced',
                        random_state=42)

```

```

rfecv = RFECV(estimator=clf, step=step, cv=StratifiedKFold(2),
              scoring=scoring)

```

```

rfecv.fit(X_train, y_train)

```

```

print("Optimal number of features : %d" % rfecv.n_features_)

```

```

# Plot number of features VS. cross-validation scores

```

```

plt.figure()

```

```

plt.xlabel(f"Number of recursive steps of {100*step:0.0f}% through {X_train.shape[1]} v

```

```

plt.ylabel(f"Cross validation score \n({scoring})")

```

```

plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)

```

```

plt.show()

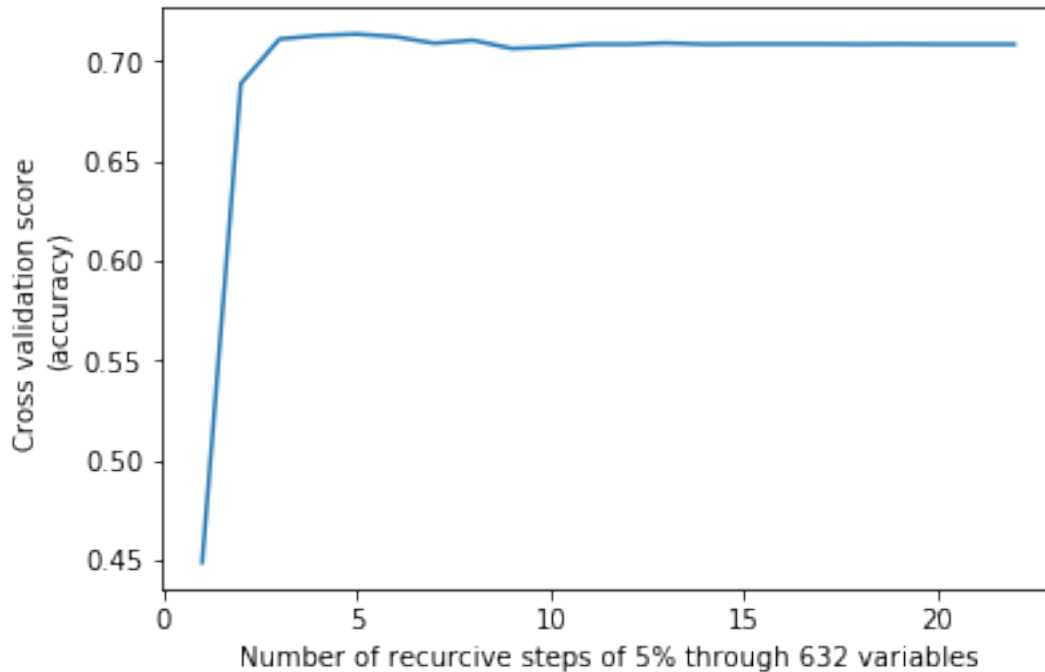
```

```

print(f"Détermination des features optimales en %0.1f s" % (time.time() - startTime))

```

Optimal number of features : 105



Détermination des features optimales en 314.2 s

1.4.2 b) Using selected features to fit various models

```
In [26]: lasso_mask = rfecv.support_.copy()
X_train = X_train[:,lasso_mask]
X_test = X_test[:,lasso_mask]
print(f"Number of features: p={X_train.shape[1]}")
```

Number of features: p=105

Logistic regression using L1 and L2

```
In [27]: nb_value = 20 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
mean_score_l2 = np.zeros(nb_value)
C_log = np.logspace(-2.5,2,nb_value)
cv = 6 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
std_scores_l1 = np.empty(nb_value)

mean_score_l2 = np.empty(nb_value)
```

```

std_scores_l2 = np.empty(nb_value)

np.random.seed(seed=42)

startTime = time.time()

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1',
                             tol=0.01, random_state=42,
                             class_weight='balanced')
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight='balanced')
    mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

    std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:] - 0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:] + 0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:] - 0.5*std_scores_l2[:], 'g--', label=u'+/-0.5 écart t')
plt.semilogx(C_log,mean_score_l2[:] + 0.5*std_scores_l2[:],'g--')


plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C \n par validation croisée \n")

```

```

(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))
print("Pénalisation l2, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l2)]))

```



```

Détermination des paramètres optimaux en 109.1 s
Pénalisation l1, valeur optimale : C = 0.74
Pénalisation l2, valeur optimale : C = 0.02

```

```

In [28]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l1',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)

```

```

print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Model score

- Accuracy : 72.0 %
- Precision : 80.4 % (Happy # positive class)
- Recall : 75.3 %
- F1 score : 77.7 %

```

In [29]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmax(mean_score_l2)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Model score

- Accuracy : 72.2 %
- Precision : 80.6 % (Happy # positive class)
- Recall : 75.2 %
- F1 score : 77.8 %

Random Forest

```

In [30]: startTime = time.time()
n_estimators_range = [16,32,64,128]
max_depth_range = [2,4,8,16,32,64,128,256]
param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)

params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced'}
clf = RandomForestClassifier(**params)

```

```

grid = GridSearchCV(clf, scoring='accuracy', param_grid=param_grid)
grid.fit(X_train, y_train)
print(f"Determination of optimal hyperparameters in {time.time() - startTime:0.1f} s")
print(f"Optimal values are {grid.best_params_} \n\
Accuracy Score of cross validation {100*grid.best_score_:0.2f}%")

# Learning on full training set with optimals hyperparameters and score on test set
params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced',
          'n_estimators' : grid.best_params_['n_estimators'],
          'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params).fit(X_train, y_train)
clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)

print(f"Random Forest, p={X_train.shape[1]}")
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

```

Determination of optimal hyperparameters in 75.3 s
Optimal values are {'max_depth': 32, 'n_estimators': 128}
Accuracy Score of cross validation 73.93%
Random Forest, p=105
Model score
- Accuracy : 73.5 %
- Precision : 74.7 % (Happy # positive class)
- Recall : 89.5 %
- F1 score : 81.4 %

```

1.4.3 b) Finding optimal set of features of given size by recursive elimination using "Lasso" (RFE)

```

In [31]: # subsetting dataframe
features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
pred = "HEUREUX_CLF"

# treating remaining missing values
print(f"{df.shape[0]} exemples before dropping na")
df_tmp = df.loc[:,set(features) | {pred}].dropna()

```

```

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42
                                                    )

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

11131 exemples before dropping na

Number exemple: 10596

- training set: 8476

- test set: 2120

Number of features: p=632

Number of class: 2

class 0 : 34.9%

class 1 : 65.1%

```

In [32]: startTime = time.time()
n_features_to_select = 20
step = 0.05
clf = LogisticRegression(C=1,
                        penalty='l1',
                        class_weight='balanced',
                        random_state=42)

selector = RFE(estimator=clf, n_features_to_select=n_features_to_select, step=step)
selector.fit(X_train, y_train)
print(f"Optimal support of size {n_features_to_select} found in {time.time() - startTime} s

```

Optimal support of size 20 found in 208.8 s

```

In [33]: small_mask = selector.support_.copy()
X_train = X_train[:,small_mask]
X_test = X_test[:,small_mask]
print(f"Number of features: p={X_train.shape[1]}")

```

Number of features: p=20


```
In [34]: print(f"Selected {X_train.shape[1]} features:\n")
        print(", ".join(X.columns[small_mask]))
```

Selected 20 features:

ACM8, NIVFRAN, NOT_AMIS, NOT_FAMI, revtot7, ETATSAN, ACM5, NIVPERSO, NIVPERS4, JUSTICE, ACM4, AC

1.4.4 c) Using selected features to fit various models

Logistic regression using L1 and L2

```
In [35]: nb_value = 20 # Nombre de valeurs testées pour l'hyperparamètre
        mean_score_l1 = np.zeros(nb_value)
        mean_score_l2 = np.zeros(nb_value)
        C_log = np.logspace(-2.5, 2, nb_value)
        cv = 6 # V-fold, nombre de fold

        mean_score_l1 = np.empty(nb_value)
        std_scores_l1 = np.empty(nb_value)

        mean_score_l2 = np.empty(nb_value)
        std_scores_l2 = np.empty(nb_value)

        np.random.seed(seed=42)

        startTime = time.time()

        for i, C in enumerate(C_log):
            clf = LogisticRegression(C=C, penalty='l1',
                                    tol=0.01, random_state=42,
                                    class_weight='balanced')
            mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                            X_train,
                                                            y_train,
                                                            cv=cv,
                                                            scoring='accuracy'))
            std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                            X_train,
                                                            y_train,
                                                            cv=cv,
                                                            scoring='accuracy'))

        for i, C in enumerate(C_log):
            clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight='balanced')
            mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                            X_train,
```

```

y_train,
cv=cv,
scoring='accuracy'))

std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
X_train,
y_train,
cv=cv,
scoring='accuracy'))

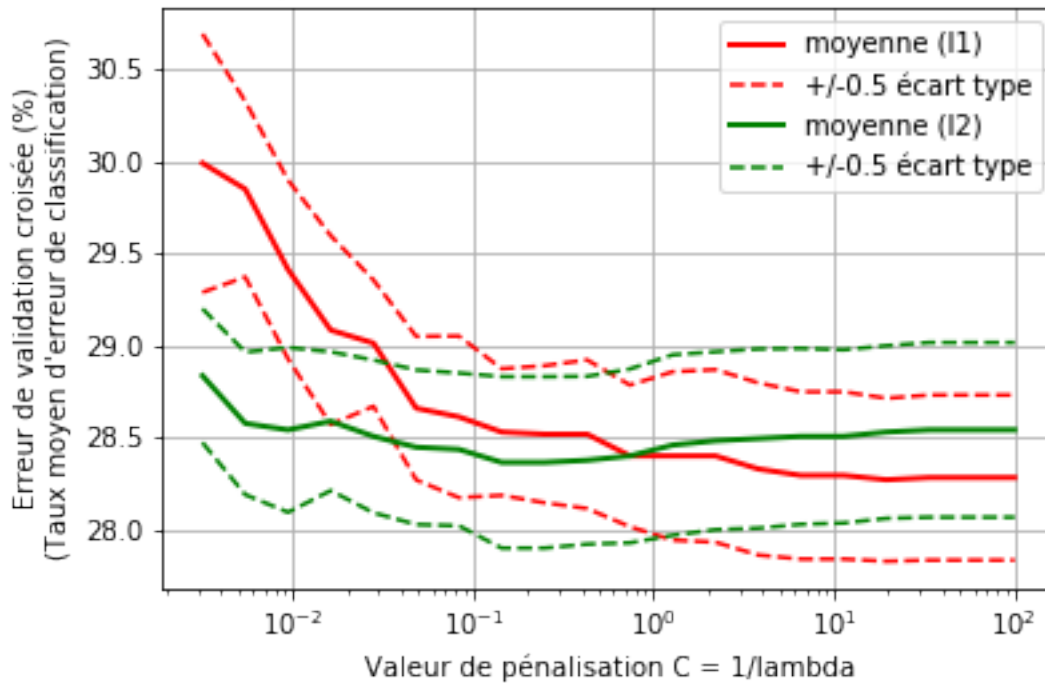
plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:] - 0.5*std_scores_l1[:],
'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:] + 0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:] - 0.5*std_scores_l2[:], 'g--', label=u'+/-0.5 écart t
plt.semilogx(C_log,mean_score_l2[:] + 0.5*std_scores_l2[:],'g--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C \npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))
print("Pénalisation l2, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l2)]))

```

Choix de l'hyperparamètre C par validation croisée (V-fold avec V = 6)



Détermination des paramètres optimaux en 19.3 s

Pénalisation l1, valeur optimale : $C = 19.47$

Pénalisation l2, valeur optimale : $C = 0.25$

```
In [36]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l1',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
```

Model score

- Accuracy : 70.6 %

- Precision : 79.4 % (Happy # positive class)
- Recall : 73.7 %
- F1 score : 76.5 %

```
In [37]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l2)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
```

Model score

- Accuracy : 70.5 %
- Precision : 79.3 % (Happy # positive class)
- Recall : 73.7 %
- F1 score : 76.4 %

```
In [38]: # Use regression coefficients to rank features
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l2)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
coef_l2 = abs(clf.coef_)
coef_sorted_l2 = -np.sort(-coef_l2).reshape(-1)
print(coef_sorted_l2)
features_sorted_l2 = np.argsort(-coef_l2).reshape(-1)
print(features_sorted_l2)
features_name = np.array(features)
features_name_sorted_l2 = features_name[features_sorted_l2]

clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l2',
                        random_state=42,
```

```

class_weight='balanced')

clf.fit(X_train,y_train)
coef_l1 = abs(clf.coef_)
coef_sorted_l1 = -np.sort(-coef_l1).reshape(-1)
features_sorded_l1 = np.argsort(-coef_l1).reshape(-1)
features_name_sorted_l1 = features_name[features_sorded_l1]

nf = X_train.shape[1]

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
ind = np.arange(nf)    # the x locations for the groups

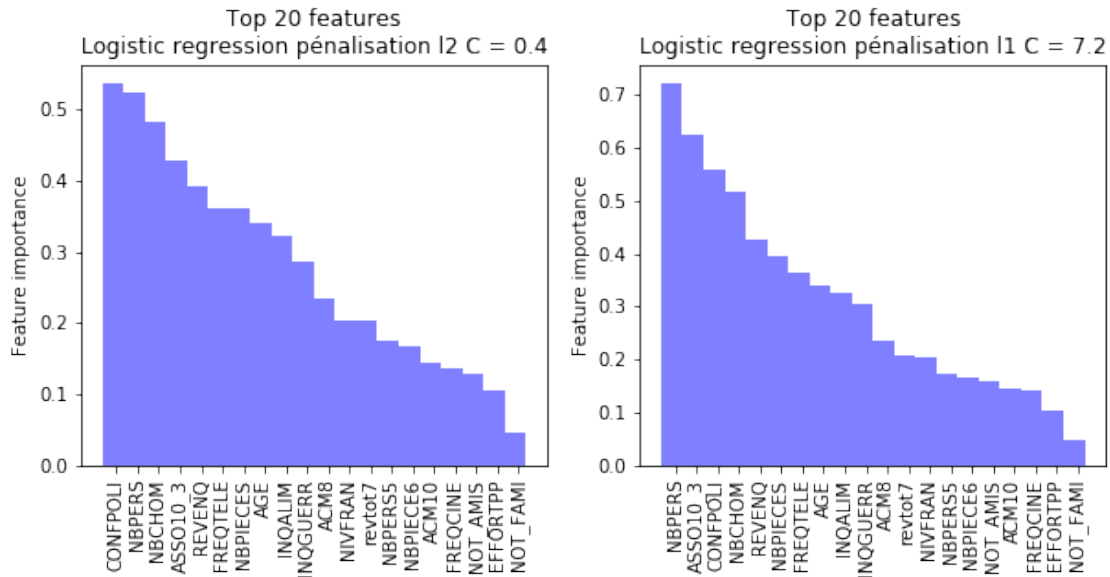
plt.subplot(1, 2, 1)
p1 = plt.bar(ind, coef_sorted_l2[0:nf], 1, color='b',alpha=0.5)
plt.ylabel('Feature importance')
plt.title(u'Top %i features\nLogistic regression pénalisation l2 C = 0.4' % nf)
plt.xticks(ind + 0.35/2.0, features_name_sorted_l2[0:nf], rotation = 90)

plt.subplot(1, 2, 2)
p1 = plt.bar(ind, coef_sorted_l1[0:nf], 1, color='b',alpha=0.5)
plt.ylabel('Feature importance')
plt.title(u'Top %i features\nLogistic regression pénalisation l1 C = 7.2' % nf)
plt.xticks(ind + 0.35/2.0, features_name_sorted_l1[0:nf], rotation = 90)

plt.show()

[ 0.53592704  0.52387649  0.48346196  0.42868429  0.39190331  0.3607874
 0.36007564  0.33949143  0.32260126  0.28526395  0.23551806  0.20346954
 0.20276202  0.17451806  0.16583378  0.14375142  0.13709309  0.12813629
 0.10392754  0.04596072]
[ 7  1  6 11  0 12  9 13 10  8  2  3 17 18 15 16 19  5  4 14]

```



Random Forest

```
In [39]: startTime = time.time()
n_estimators_range = [32,64,128,256,512]
max_depth_range = [4,8,16,32,64]
param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)

params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced'}
clf = RandomForestClassifier(**params)

grid = GridSearchCV(clf, scoring='accuracy', param_grid=param_grid)
grid.fit(X_train, y_train)
print(f"Determination of optimal hyperparameters in {time.time() - startTime:0.1f} s")
print(f"Optimal values are {grid.best_params_} \n\
Accuracy Score of cross validation {100*grid.best_score_:0.2f}%")

# Learning on full training set with optimals hyperparameters and score on test set
params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced',
          'n_estimators' : grid.best_params_['n_estimators'],
          'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params).fit(X_train, y_train)
clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)

print(f"Random Forest, p={X_train.shape[1]}")
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
```

```

p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Determination of optimal hyperparameters in 141.7 s
 Optimal values are {'max_depth': 16, 'n_estimators': 512}
 Accuracy Score of cross validation 71.79%
 Random Forest, p=20
 Model score
 - Accuracy : 71.7 %
 - Precision : 75.0 % (Happy # positive class)
 - Recall : 84.7 %
 - F1 score : 79.5 %

1.5 5) Multi class regression

1.5.1 a) Training and test set preparation

```

In [40]: features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
        pred = "HEUREUX"

        # treating remaining missing values
        df_tmp = df.loc[:,set(features) | {pred} ]
        df_tmp.loc[df_tmp["HEUREUX"]==5, "HEUREUX"]=None
        df_tmp = df_tmp.dropna()

        X = df_tmp.loc[:,features]
        y = df_tmp[pred]

        X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                            test_size=0.2,
                                                            random_state=42)

        scaler = StandardScaler().fit(X_train)
        X_train = scaler.transform(X_train)
        X_test = scaler.transform(X_test)

        print(f"Number exemple: {y.shape[0]}\n- training set: \
        {y_train.shape[0]}\n- test set: {y_test.shape[0]}")
        print(f"Number of features: p={X_train.shape[1]}")
        print(f"Number of class: {len(np.unique(y))}")
        for c in np.unique(y):
            print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

Number exemple: 10596

```

- training set: 8476
- test set: 2120
Number of features: p=632
Number of class: 4
class 1 : 1.8%
class 2 : 33.1%
class 3 : 49.1%
class 4 : 15.9%

```

1.5.2 b) Feature selection

```
In [41]: startTime = time.time()
```

```

scoring='f1_macro'
step = 0.05

```

```

clf = LogisticRegression(C=1,
                        penalty='l1',
                        class_weight='balanced',
                        random_state=42)

```

```

rfecv = RFECV(estimator=clf, step=step, cv=StratifiedKFold(2),
              scoring=scoring)

```

```
rfecv.fit(X_train, y_train)
```

```
print("Optimal number of features : %d" % rfecv.n_features_)
```

```
# Plot number of features VS. cross-validation scores
```

```
plt.figure()
```

```
plt.xlabel(f"Number of recursive steps of {100*step:0.0f}% through {X_train.shape[1]} v
```

```
plt.ylabel(f"Cross validation score \n({scoring})")
```

```
plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)
```

```
plt.show()
```

```
print(f"Détermination des features optimales en %0.1f s" % (time.time() - startTime))
```

```

//anaconda/envs/py36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: Undefined
'precision', 'predicted', average, warn_for)
//anaconda/envs/py36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: Undefined
'precision', 'predicted', average, warn_for)

```

```
Optimal number of features : 198
```




Détermination des features optimales en 4920.4 s

```
In [42]: nclf_mask = rfecv.support_.copy()
X_train = X_train[:,nclf_mask]
X_test = X_test[:,nclf_mask]
print(f"Number of features: p={X_train.shape[1]}")
```

Number of features: p=198

```
In [43]: nb_value = 10 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
mean_score_l2 = np.zeros(nb_value)
C_log = np.logspace(-2,2,nb_value)
cv = 3 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
std_scores_l1 = np.empty(nb_value)

mean_score_l2 = np.empty(nb_value)
std_scores_l2 = np.empty(nb_value)

np.random.seed(seed=42)
```

```

startTime = time.time()

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1', tol=0.01, random_state=42, class_weight=
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight=
    mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

    std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))


plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:] - 0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:] + 0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:] - 0.5*std_scores_l2[:], 'g--',
             label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l2[:] + 0.5*std_scores_l2[:],'g--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel("F1 macro de validation croisée (%)")
plt.title(u"Choix de l'hyperparamètre C\npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print(f"Détermination des paramètres optimaux en \

```

```
{time.time() - startTime:0.1f} s")
print(f"Pénalisation l1, valeur optimale : \
C = {C_log[np.argmax(mean_score_l1)]:0.4f}")
print(f"Pénalisation l2, valeur optimale : \
C = {C_log[np.argmax(mean_score_l2)]:0.4f}")
```



Détermination des paramètres optimaux en 3646.7 s

Pénalisation l1, valeur optimale : C = 0.0100

Pénalisation l2, valeur optimale : C = 0.0100

```
In [44]: # Learning on full training set with optimals hyperparameters and score on test set
clf = LogisticRegression(C=C_log[np.argmax(mean_score_l2)],
                        penalty='l2',
                        tol=0.01,
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
```

```
In [45]: def plot_confusion_matrix(cm, classes,
                                normalize=False,
```

```

        title='Confusion matrix',
        cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

```

```

In [46]: # Model evaluation
class_names = ["Jamais",
               "Occasionnellement",
               "Assez souvent",
               "Très souvent" ]

f1_scores = f1_score(y_test, y_test_pred, labels = [1,2,3,4], average=None)
for i,c in enumerate(class_names):
    print(f"f1 score class '{c}' : {100*f1_scores[i]:0.1f}%")
accuracy = clf.score(X_test, y_test)
f1_macro = f1_score(y_test, y_test_pred, average='macro')
f1_weighted = f1_score(y_test, y_test_pred, average='weighted')
print(f"Average scores :\nf1 macro : {f1_macro*100:0.4f} %\n\
f1 weighted : {f1_weighted*100:0.4f} %\naccuracy : {accuracy*100:0.4f} %")

# Compute confusion matrix
cnf_matrix = confusion_matrix(y_test, y_test_pred)

```

```

np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()

f1 score class 'Jamais' : 16.1%
f1 score class 'Occasionnellement' : 53.4%
f1 score class 'Assez souvent' : 60.1%
f1 score class 'Très souvent' : 43.4%
Average scores :
f1 macro : 43.2525 %
f1 weighted : 54.1814 %
accuracy : 53.4434 %
Confusion matrix, without normalization
[[ 12  26   0   5]
 [ 54 354 236  58]
 [ 30 208 601 180]
 [ 10  35 145 166]]
Normalized confusion matrix
[[ 0.28  0.6   0.   0.12]
 [ 0.08  0.5   0.34  0.08]
 [ 0.03  0.2   0.59  0.18]
 [ 0.03  0.1   0.41  0.47]]

```

