

# Felix\_prototype\_V0

October 27, 2018

## 1 Felix prototype

Version 0

Date 21/10/2018

Model used : **Random Forest Classifier** on features selected through **lasso**  
Clustering method used : **Hierarchical clustering** using **ward metric** based on 6 **NOT variable**

```
In [1]: from pathlib import Path
import pandas as pd
import numpy as np
from datetime import datetime
import time
import matplotlib.pyplot as plt
%matplotlib inline
import itertools
import pickle
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, f1_score, precision_score, recall_score
from sklearn.model_selection import StratifiedKFold
from sklearn.utils import resample

In [2]: path_project = Path.home() / Path('Google Drive/Felix')
path_data = path_project / Path("data")
path_dump = path_project / Path("dump")

In [3]: # loading data
file = path_data / Path("dataset.csv")
with Path.open(file, 'rb') as fp:
    dataset = pd.read_csv(fp, encoding='utf-8', low_memory=False, index_col = 0)
```

### 1.0.1 Features scope and selection strategy

Features are selected using lasso on the full scope of feature. The 50 more important features (logistic regression coef ranking) are kept regardless of their activability



```

df.loc[df["HEUREUX"]==3, "HEUREUX_CLF"] = 1
df.loc[df["HEUREUX"]==5, "HEUREUX_CLF"] = None

scope = ( RFE_LogisticRegression_10_features ) & set(dataset.columns)
n_max = 2000

df = df.loc[:,scope | {"HEUREUX_CLF"} ].dropna()
features = df.loc[:,scope ].columns

X = df.loc[:,scope]
y = df["HEUREUX_CLF"]

Xs, ys = resample(X, y, random_state=42)

Xs = Xs.iloc[0:n_max,:]
ys = ys.iloc[0:n_max]

X_train, X_test, y_train, y_test = train_test_split(Xs, ys,
                                                    test_size=0.2,
                                                    random_state=42
                                                    )

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

```

Number exemple: 10915
- training set: 1600
- test set: 400
Number of features: p=10
Number of class: 2
class 0 : 34.9%
class 1 : 65.1%

```

#### 1.0.4 Learning and model performance evaluation on full dataset (before clustering)

```

In [9]: startTime = time.time()
        n_estimators_range = [32,64,128,256,512]
        max_depth_range = [4,8,16,32,64]

```

```

param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)

params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced'}
clf = RandomForestClassifier(**params)

grid = GridSearchCV(clf, scoring='accuracy', param_grid=param_grid)
grid.fit(X_train, y_train)
print(f"Determination of optimal hyperparameters in {time.time() - startTime:0.1f} s")
print(f"Optimal values are {grid.best_params_} \n\
Accuracy Score of cross validation {100*grid.best_score_:0.2f}%")

# Learning on full training set with optimal hyperparameters and score on test set
params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced',
          'n_estimators' : grid.best_params_['n_estimators'],
          'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params).fit(X_train, y_train)
clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)

print(f"Random Forest, p={X_train.shape[1]}")
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
res_full = {
    'f1_score' : f1,
    'accuracy' : accuracy,
    'precision' : p,
    'recall' : r
}

```

```

Determination of optimal hyperparameters in 34.9 s
Optimal values are {'max_depth': 8, 'n_estimators': 64}
Accuracy Score of cross validation 71.44%
Random Forest, p=10
Model score
- Accuracy : 70.2 %
- Precision : 78.3 % (Happy # positive class)
- Recall : 74.3 %
- F1 score : 76.2 %

```

```
In [10]: importances = clf.feature_importances_
```

```

std = np.std([tree.feature_importances_ for tree in clf.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]
features_name = np.array(features)
#features_name_sorted_rf = features_name[indices]
# Print the feature ranking
print("Feature ranking:")

max_features = 15

actionable_individual_1_features = cdv_actionable_individual_1_features
actionable_individual_2_features = cdv_actionable_individual_2_features
actionable_admin_1_features = cdv_actionable_admin_1_features | insee_recreation_action
actionable_admin_2_features = cdv_actionable_admin_2_features | insee_recreation_action

for f in range(min(X.shape[1],max_features)):
    print("%d. feature %d -%s- (%f)" % (f + 1, indices[f], features_name[indices[f]], im
    if features_name[indices[f]] in actionable_individual_1_features:
        print("\tActionable at individual level (1)")
    if features_name[indices[f]] in actionable_individual_2_features:
        print("\tActionable at individual level (2)")
    if features_name[indices[f]] in actionable_admin_1_features:
        print("\tActionable at administrative level (1)")
    if features_name[indices[f]] in actionable_admin_2_features:
        print("\tActionable at administrative level (2)")


# Plot the feature importances of the forest
plt.figure()
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X.shape[1]), indices)
plt.xlim([-1, X.shape[1]])
plt.show()

```

Feature ranking:

1. feature 5 -NIVPERSO- (0.158025)  
     Actionable at individual level (2)  
     Actionable at administrative level (2)
2. feature 7 -NOT\_AMIS- (0.144404)  
     Actionable at individual level (1)  
     Actionable at administrative level (2)
3. feature 6 -CADVIE- (0.138970)  
     Actionable at individual level (1)

- Actionable at administrative level (1)
- 4. feature 0 -NOT\_FAMI- (0.124023)
  - Actionable at individual level (1)
  - Actionable at administrative level (2)
- 5. feature 4 -ETATSAN- (0.103254)
  - Actionable at individual level (1)
  - Actionable at administrative level (1)
- 6. feature 9 -SOUFFDEP\_Oui- (0.084266)
  - Actionable at individual level (2)
  - Actionable at administrative level (1)
- 7. feature 2 -RE\_ALIM\_Oui- (0.077853)
  - Actionable at individual level (2)
  - Actionable at administrative level (2)
- 8. feature 3 -LIEN\_2\_Conjoint ou compagnon- (0.069062)
- 9. feature 8 -SOUFFNER\_Oui- (0.052739)
  - Actionable at individual level (2)
  - Actionable at administrative level (1)
- 10. feature 1 -SITUEMP3\_Inactif- (0.047405)
  - Actionable at individual level (2)
  - Actionable at administrative level (2)



### 1.0.5 Learning and model performance evaluation on each clusters

```
In [11]: n_estimators_range = [16,32,64,128]
max_depth_range = [2,4,8,16,32,64]
param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)
params = {'max_features' : 'sqrt',
          'random_state' : 32,
          'min_samples_split' : 2,
          'class_weight' : 'balanced'
        }

scope = ( RFE_LogisticRegression_50_features ) & set(dataset.columns)
features = df.loc[:,scope].columns

In [12]: score_clustering_methods = []
clustering_methods = clustTest1.columns[2:3]

for method in clustering_methods:
    print("-----")
    print(f"\nAnalysis cluster method {method}")
    cluster_list = clustTest1[method].unique()
    print(f"liste of clusters : {cluster_list}")
    score_cluster = []
    for cluster in cluster_list:
        index_scope = clustTest1.loc[clustTest1[method]==cluster,:].index
        print(f"cluster {cluster} : {len(index_scope)} elements")

        Xc = X.loc[index_scope.intersection(X.index),:]
        yc = y[index_scope.intersection(X.index)]

        Xs, ys = resample(Xc, yc, random_state=42)

        Xs = Xs.iloc[0:n_max,:]
        ys = ys.iloc[0:n_max]

        X_train, X_test, y_train, y_test = train_test_split(Xs, ys,
                                                            test_size=0.2,
                                                            random_state=42)

        scaler = StandardScaler().fit(X_train)
        X_train = scaler.transform(X_train)
        X_test = scaler.transform(X_test)

    print(f"Number exemple: {ys.shape[0]}\n\
- training set: {y_train.shape[0]}\n\
- test set: {y_test.shape[0]}")
    print(f"Number of features: p={X_train.shape[1]}")
    print(f"Number of class: {len(np.unique(y))}")
    for c in np.unique(y):
        print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")
```

```

startTime = time.time()
clf = RandomForestClassifier(**params)
grid = GridSearchCV(clf,
                    scoring='accuracy',
                    param_grid=param_grid)

grid.fit(X_train, y_train)
print(f"Optimal values are {grid.best_params_} \n\
cross validation score {100*grid.best_score_:0.2f}%")
print()

# Learning on full training set with optimal hyperparameters and score on test
params_opt = {'max_features' : 'sqrt', 'random_state' : 32,
              'min_samples_split' : 2, 'class_weight' : 'balanced',
              'n_estimators' : grid.best_params_['n_estimators'],
              'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params_opt).fit(X_train, y_train)

y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)

res = {'f1_score' : f1,
       'accuracy' : accuracy,
       'precision' : p,
       'recall' : r}

cl = {'cluster' : cluster,
      'size' : len(index_scope),
      'model' : 'RandomForestClassifier',
      'params' : params_opt,
      'metrics' : res
      }

score_cluster.append(cl)

d = {'clustering_method' : method,
     'cluster_scores' : score_cluster
     }
score_clustering_methods.append(d)

```

---



Analysis cluster method clust3  
 liste of clusters : [2 4 6 1 3 5]  
 cluster 2 : 3053 elements  
 Number exemple: 2000  
     - training set: 1600  
     - test set: 400  
 Number of features: p=10  
 Number of class: 2  
 class 0 : 34.9%  
 class 1 : 65.1%  
 Optimal values are {'max\_depth': 16, 'n\_estimators': 64}  
 cross validation score 74.56%

cluster 4 : 2359 elements  
 Number exemple: 2000  
     - training set: 1600  
     - test set: 400  
 Number of features: p=10  
 Number of class: 2  
 class 0 : 34.9%  
 class 1 : 65.1%  
 Optimal values are {'max\_depth': 16, 'n\_estimators': 64}  
 cross validation score 80.38%

cluster 6 : 2313 elements  
 Number exemple: 2000  
     - training set: 1600  
     - test set: 400  
 Number of features: p=10  
 Number of class: 2  
 class 0 : 34.9%  
 class 1 : 65.1%  
 Optimal values are {'max\_depth': 16, 'n\_estimators': 64}  
 cross validation score 77.69%

cluster 1 : 528 elements  
 Number exemple: 521  
     - training set: 416  
     - test set: 105  
 Number of features: p=10  
 Number of class: 2  
 class 0 : 34.9%  
 class 1 : 65.1%  
 Optimal values are {'max\_depth': 32, 'n\_estimators': 128}  
 cross validation score 77.16%

cluster 3 : 1384 elements  
 Number exemple: 1374

```

- training set: 1099
- test set: 275
Number of features: p=10
Number of class: 2
class 0 : 34.9%
class 1 : 65.1%
Optimal values are {'max_depth': 32, 'n_estimators': 64}
cross validation score 82.07%

```

```

cluster 5 : 1494 elements
Number exemple: 1489
- training set: 1191
- test set: 298
Number of features: p=10
Number of class: 2
class 0 : 34.9%
class 1 : 65.1%
Optimal values are {'max_depth': 16, 'n_estimators': 64}
cross validation score 79.51%

```

## 1.0.6 Performance gain obtained using clustering

```

In [13]: # F1 score
for score_method in score_clustering_methods:
    print(f"method {score_method['clustering_method']}:")
    average_score = 0
    total_size = 0
    for i, score_cluster in enumerate(score_method['cluster_scores']):
        print(f"cluster {score_cluster['cluster']} ({score_cluster['size']}), f1 macro {score_cluster['f1_macro']}")
        average_score += score_cluster['metrics']['f1_score']*score_cluster['size']
        total_size += score_cluster['size']

    average_score = average_score / total_size
    print(f"average f1 on clusters {100*average_score:0.1f}% gain {100*(average_score-r

method clust3:
cluster 2 (3053), f1 macro 83.8%
cluster 4 (2359), f1 macro 86.9%
cluster 6 (2313), f1 macro 85.0%
cluster 1 (528), f1 macro 74.7%
cluster 3 (1384), f1 macro 87.1%
cluster 5 (1494), f1 macro 87.1%
average f1 on clusters 85.1% gain 8.9

```

```

In [14]: # accuracy

```

```

for score_method in score_clustering_methods:
    print(f"method {score_method['clustering_method']}:")
    average_score = 0
    total_size = 0
    for i, score_cluster in enumerate(score_method['cluster_scores']):
        print(f"cluster {score_cluster['cluster']} ({score_cluster['size']}) , accuracy {score_cluster['accuracy']}")
        average_score = average_score + score_cluster['metrics']['accuracy']*score_cluster['size']
        total_size += score_cluster['size']
    average_score = average_score / total_size
    print(f"average accuracy on clusters {100*average_score:0.1f}% gain {100*(average_score - 0.8):0.1f}%")

method clust3:
cluster 2 (3053) , accuracy 78.2%
cluster 4 (2359) , accuracy 82.5%
cluster 6 (2313) , accuracy 78.5%
cluster 1 (528) , accuracy 76.2%
cluster 3 (1384) , accuracy 82.9%
cluster 5 (1494) , accuracy 84.2%
average accuracy on clusters 80.5% gain 10.2

```

### 1.0.7 Feature importance of the models & actionable variables