

# Models using full features scope

September 2, 2018

## 1 Firts attempts & Base Model

```
In [1]: from pathlib import Path
import pandas as pd
import numpy as np
from datetime import datetime
import time
import matplotlib.pyplot as plt
%matplotlib inline
#%pylab inline
import itertools
import pickle
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.decomposition import PCA
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, f1_score, precision_score, recall_score
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import LabelBinarizer
from sklearn.preprocessing import OneHotEncoder
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from sklearn.feature_selection import RFECV, RFE

In [2]: path_project = Path.home() / Path('Google Drive/Felix')
path_data = path_project / Path("data")
path_dump = path_project / Path("dump")

In [3]: # loading cdv data
file = path_data / Path("felix.csv")
with Path.open(file, 'rb') as fp:
    cdv = pd.read_csv(fp, encoding='cp1252', low_memory=False, index_col = 0)

In [4]: # loadind cdv data without format
file = path_data / Path("felix_ssfmt.csv")
```

```

with Path.open(file, 'rb') as fp:
    cdv_ssfmt = pd.read_csv(fp, encoding='cp1252', low_memory=False, index_col = 0)

```

```

In [5]: # loading MergeCommunesEnvi data
file = path_data / Path("MergeCommunesEnvi.csv")
with Path.open(file, 'rb') as fp:
    MergeCommunesEnvi = pd.read_csv(fp, encoding='cp1252', low_memory=False, sep=';', index_col = 0)

```

## 1.1 1) Feature engineering

```

In [6]: filename = path_dump / Path("dict_var_groups.sav")
with open(filename, 'rb') as fp:
    dict_var_groups = pickle.load(fp)

scope_2015_var = dict_var_groups['scope_2015_var']
scope_2016_var = dict_var_groups['scope_2016_var']
scope_2017_var = dict_var_groups['scope_2017_var']
scope_2018_var = dict_var_groups['scope_2018_var']
scope_2015_2018_var = dict_var_groups['scope_2015_2018_var']
scope_2016_2018_var = dict_var_groups['scope_2016_2018_var']
scope_2017_2018_var = dict_var_groups['scope_2017_2018_var']
pred_var = dict_var_groups['pred_var']
com_var = dict_var_groups['com_var']
tech_var = dict_var_groups['tech_var']
text_var = dict_var_groups['text_var']
bizz_var = dict_var_groups['bizz_var']
cat_var = dict_var_groups['cat_var']
cat_max9_var = dict_var_groups['cat_max9_var']
cat_min10_var = dict_var_groups['cat_min10_var']
quant_var = dict_var_groups['quant_var']

indiv_semi_act_var = dict_var_groups["indiv_semi_act_var"]
indiv_act_var = dict_var_groups["indiv_act_var"]
admin_semi_act_var = dict_var_groups["admin_semi_act_var"]
admin_act_var = dict_var_groups["admin_act_var"]
commune_var = dict_var_groups["commune_var"]

scope_2015_ext_var = dict_var_groups["scope_2015_ext_var"]
scope_2016_ext_var = dict_var_groups["scope_2016_ext_var"]
scope_2017_ext_var = dict_var_groups["scope_2017_ext_var"]
scope_2018_ext_var = dict_var_groups["scope_2018_ext_var"]
scope_2015_2018_ext_var = dict_var_groups["scope_2015_2018_ext_var"]
scope_2016_2018_ext_var = dict_var_groups["scope_2016_2018_ext_var"]
scope_2017_2018_ext_var = dict_var_groups["scope_2017_2018_ext_var"]

In [7]: exclusion = com_var | tech_var | bizz_var | text_var
scope_kept = scope_2015_2018_ext_var - exclusion
print(f"{len(scope_kept)} are kept ")

```

```

cat_var_kept = cat_max9_var & scope_kept
scope_quant_var = (quant_var & scope_kept)
quant_null = np.sum(MergeCommunesEnvi.loc[:,scope_quant_var].isnull())
quant_var_kept = set(quant_null[quant_null < 200].index)
print(f"out of the {len(scope_quant_var)} quantitatives variables:")
print(f"{len(quant_var_kept)} have less than 200 missing values and are kept")

```

464 are kept  
out of the 290 quantitatives variables:  
263 have less than 200 missing values and are kept

```

In [8]: scope = cat_var_kept | quant_var_kept
df = MergeCommunesEnvi.loc[:,scope]
# ...
df.loc[:,scope_2015_2018_var & scope] = cdv_ssfmt.loc[:,scope_2015_2018_var & scope]
df.loc[:,cat_var_kept - {"HEUREUX"}] = cdv.loc[:,cat_var_kept - {"HEUREUX"}]

```

```

In [9]: print(f"Number of variable kept {df.shape[1]}")

```

Number of variable kept 419

### 1.1.1 Encoding 'HEUREUX'

```

In [10]: # reducing problem to a 2 class classification problem
df["HEUREUX_CLF"] = 0
df.loc[df["HEUREUX"]==4, "HEUREUX_CLF"] = 1
df.loc[df["HEUREUX"]==3, "HEUREUX_CLF"] = 1
df.loc[df["HEUREUX"]==5, "HEUREUX_CLF"] = None

```

```

In [11]: # Modelisation as a regression problem
df["HEUREUX_REG"] = df["HEUREUX"]
df.loc[df["HEUREUX"]==5, "HEUREUX_REG"] = None

```

### 1.1.2 Encoding categorial variables

```

In [12]: p = df.shape[1]
print(f"{p} columns")
print(f"out of which {len(cat_var_kept)-1} are corresponding to categorial features")

```

421 columns  
out of which 155 are corresponding to categorial features

```

In [13]: df = pd.get_dummies(df,
                             columns=cat_var_kept - {"HEUREUX"},
                             dummy_na = True,
                             drop_first=1)

```

```
In [14]: q = df.shape[1]
         print(f"{q} columns after encoding")
         print(f"{len(cat_var_kept)-1} variables where re-encoded in {len(cat_var_kept)-1+q-p}")
```

838 columns after encoding  
155 variables where re-encoded in 572

```
In [15]: def get_related_features(variable):
         '''return all columns in global variable df
         starting by variable'''
         if isinstance(variable, str):
             scope = {variable}
         else:
             scope = set(variable)
         features = set()
         for element in scope:
             features = features | {c for c in df.columns if len(c) > len(element)
                                     and c[0:len(element)] == element
                                     and c[len(element)]=='_'}
         return features
```

## 1.2 2) Dataset construction and visualisation

```
In [16]: # subsetting dataframe
         features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
         pred = "HEUREUX_CLF"

         # treating remaining missing values
         df_tmp = df.loc[:,set(features) | {pred} ].dropna()

         X = df_tmp.loc[:,features]
         y = df_tmp[pred]


         X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                             test_size=0.2,
                                                             random_state=42)

         scaler = StandardScaler().fit(X_train)
         X_train = scaler.transform(X_train)
         X_test = scaler.transform(X_test)

         print(f"Number exemple: {y.shape[0]}\n- training set: \
               {y_train.shape[0]}\n- test set: {y_test.shape[0]}")
         print(f"Number of features: p={X_train.shape[1]}")
         print(f"Number of class: {len(np.unique(y))}")
         for c in np.unique(y):
             print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")
```

Number exemple: 10445  
- training set: 8356  
- test set: 2089  
Number of features: p=835  
Number of class: 2  
class 0 : 35.1%  
class 1 : 64.9%

```
In [17]: # Reduction dim PCA
pca = PCA(n_components=2)
pca.fit(X_train)
X_r = pca.transform(X_train)
happy = (y_train==1)
unhappy = (y_train==0)
plt.scatter(X_r[happy,0], X_r[happy,1],
            s=80, c='blue',marker=">", label="happy")
plt.scatter(X_r[unhappy,0], X_r[unhappy,1],
            s=80, c='red',marker='+', label="unhappy")
plt.ylabel(f'Projection VP2 \
{pca.explained_variance_ratio_[1]*100:0.1f}% de variance')
plt.xlabel(f'Projection VP1 \
{pca.explained_variance_ratio_[0]*100:0.1f}% de variance')
plt.title("Projection 2D par PCA")
plt.legend(bbox_to_anchor=(0.4, 0.25))
plt.show()
```



### 1.3 3) Baseline model

Logistic regression with 4 features to predict a 2 class variable

#### 1.3.1 a) Training set and test set preparation

```
In [18]: # subsetting dataframe
features = {'ETATSAN', 'NOT_AMIS', 'DEPLOY', 'NOT_FAMI'}
pred = "HEUREUX_CLF"

# treating remaining missing values
df_tmp = df.loc[:,set(features) | {pred}].dropna()

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")
```

Number exemple: 10915  
- training set: 8732  
- test set: 2183  
Number of features: p=4  
Number of class: 2  
class 0 : 34.9%  
class 1 : 65.1%

#### 1.3.2 b) Hyperparameters tuning

```
In [19]: nb_value = 50 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
C_log = np.logspace(-2,2,nb_value)
cv = 6 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
```

```

std_scores_l1 = np.empty(nb_value)

np.random.seed(seed=42)

startTime = time.time()


for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1',
                             tol=0.01, random_state=42,
                             class_weight='balanced')
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:]-0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:]+0.5*std_scores_l1[:],'r--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C\npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))

```



Détermination des paramètres optimaux en 16.2 s  
Pénalisation l1, valeur optimale :  $C = 0.91$

### 1.3.3 c) Model training and evaluation

```
In [20]: # Learning on full training set with optimals hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l1',
                        tol=0.01,
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
```



Model score

- Accuracy : 66.6 %
- Precision : 75.9 % (Happy # positive class)
- Recall : 71.5 %
- F1 score : 73.6 %

## 1.4 4) Model selection (2 class)

### 1.4.1 a) Finding optimal set of features by recursive elimination using "Lasso" (RFECV)

```
In [21]: # subsetting dataframe
features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
pred = "HEUREUX_CLF"

# treating remaining missing values
print(f"{df.shape[0]} exemples before dropping na")
df_tmp = df.loc[:, set(features) | {pred} ].dropna()

X = df_tmp.loc[:, features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42
                                                    )

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

11131 exemples before dropping na
Number exemple: 10445
- training set: 8356
- test set: 2089
Number of features: p=835
Number of class: 2
class 0 : 35.1%
class 1 : 64.9%
```

```
In [22]: startTime = time.time()
```

```

scoring='accuracy'
step = 0.05

clf = LogisticRegression(C=1,
                        penalty='l1',
                        class_weight='balanced',
                        random_state=42)

rfecv = RFECV(estimator=clf, step=step, cv=StratifiedKFold(2),
              scoring=scoring)

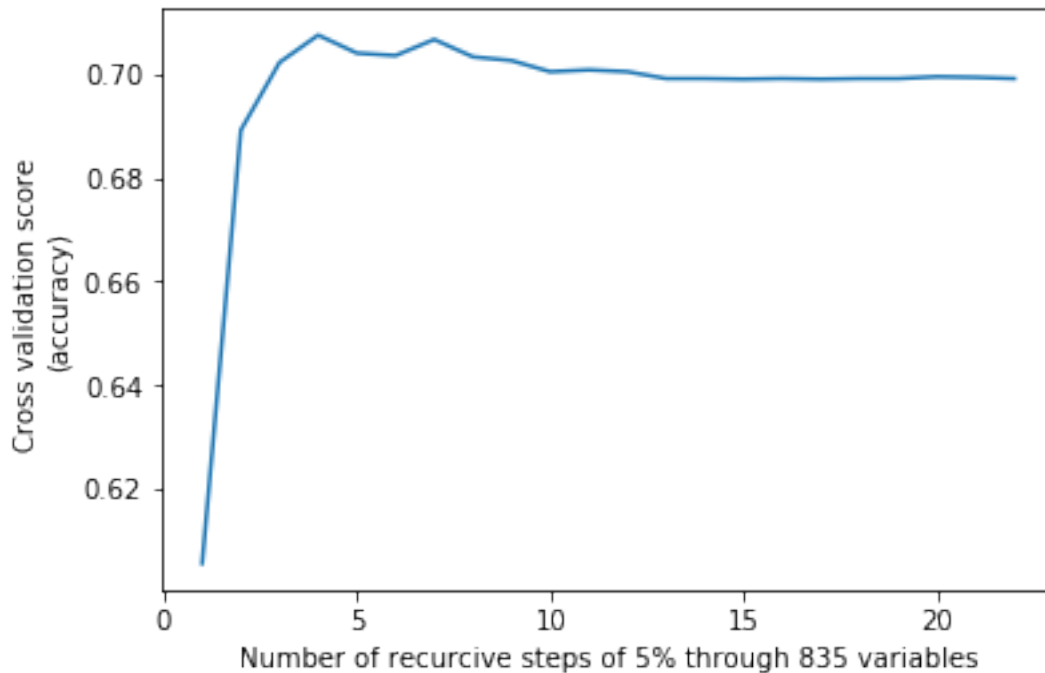
rfecv.fit(X_train, y_train)

print("Optimal number of features : %d" % rfecv.n_features_)

# Plot number of features VS. cross-validation scores
plt.figure()
plt.xlabel(f"Number of recursive steps of {100*step:0.0f}% through {X_train.shape[1]} v
plt.ylabel(f"Cross validation score \n({scoring})")
plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)
plt.show()
print(f"Détermination des features optimales en %0.1f s" % (time.time() - startTime))

```

Optimal number of features : 97



Détermination des features optimales en 1517.5 s

#### 1.4.2 b) Using selected features to fit various models

```
In [23]: lasso_mask = rfecv.support_.copy()
         X_train = X_train[:,lasso_mask]
         X_test = X_test[:,lasso_mask]
         print(f"Number of features: p={X_train.shape[1]}")
```

Number of features: p=97

#### Logistic regression using L1 and L2

```
In [24]: nb_value = 20 # Nombre de valeurs testées pour l'hyperparamètre
         mean_score_l1 = np.zeros(nb_value)
         mean_score_l2 = np.zeros(nb_value)
         C_log = np.logspace(-2.5,2,nb_value)
         cv = 6 # V-fold, nombre de fold

         mean_score_l1 = np.empty(nb_value)
         std_scores_l1 = np.empty(nb_value)

         mean_score_l2 = np.empty(nb_value)
         std_scores_l2 = np.empty(nb_value)

         np.random.seed(seed=42)

         startTime = time.time()

         for i, C in enumerate(C_log):
             clf = LogisticRegression(C=C, penalty='l1',
                                     tol=0.01, random_state=42,
                                     class_weight='balanced')
             mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                             X_train,
                                                             y_train,
                                                             cv=cv,
                                                             scoring='accuracy'))
             std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                            X_train,
                                                            y_train,
                                                            cv=cv,
                                                            scoring='accuracy'))
```

```

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight=
    mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))

    std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='accuracy'))


plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:] - 0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:] + 0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:] - 0.5*std_scores_l2[:], 'g--', label=u'+/-0.5 écart t
plt.semilogx(C_log,mean_score_l2[:] + 0.5*std_scores_l2[:],'g--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C \npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))
print("Pénalisation l2, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l2)]))

```

### Choix de l'hyperparamètre C par validation croisée (V-fold avec V = 6)



Détermination des paramètres optimaux en 181.8 s

Pénalisation l1, valeur optimale :  $C = 3.79$

Pénalisation l2, valeur optimale :  $C = 0.03$

```
In [25]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l1',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
```

Model score

- Accuracy : 70.6 %

- Precision : 80.5 % (Happy # positive class)
- Recall : 71.9 %
- F1 score : 76.0 %

```
In [26]: # Learning on full training set with optimal hyperparameters
# and score evaluation on test set
clf = LogisticRegression(C=C_log[np.argmin(mean_score_l2)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")
```

Model score

- Accuracy : 71.2 %
- Precision : 81.1 % (Happy # positive class)
- Recall : 72.4 %
- F1 score : 76.5 %

## Random Forest

```
In [27]: startTime = time.time()
n_estimators_range = [16,32,64,128]
max_depth_range = [2,4,8,16,32,64,128,256]
param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)

params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced'}
clf = RandomForestClassifier(**params)

grid = GridSearchCV(clf, scoring='accuracy', param_grid=param_grid)
grid.fit(X_train, y_train)
print(f"Determination of optimal hyperparameters in {time.time() - startTime:0.1f} s")
print(f"Optimal values are {grid.best_params_} \n\
Accuracy Score of cross validation {100*grid.best_score_:0.2f}%")

# Learning on full training set with optimal hyperparameters and score on test set
params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced',
```

```

        'n_estimators' : grid.best_params_['n_estimators'],
        'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params).fit(X_train, y_train)
clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)

print(f"Random Forest, p={X_train.shape[1]}")
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Determination of optimal hyperparameters in 69.1 s  
 Optimal values are {'max\_depth': 16, 'n\_estimators': 128}  
 Accuracy Score of cross validation 73.55%  
 Random Forest, p=97  
 Model score  
 - Accuracy : 72.5 %  
 - Precision : 74.7 % (Happy # positive class)  
 - Recall : 86.8 %  
 - F1 score : 80.3 %

### 1.4.3 b) Finding optimal set of features of given size by recursive elimination using "Lasso" (RFE)

```

In [28]: # subsetting dataframe
features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
pred = "HEUREUX_CLF"

# treating remaining missing values
print(f"{df.shape[0]} exemples before dropping na")
df_tmp = df.loc[:,set(features) | {pred}].dropna()

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42
                                                    )

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)

```

```

X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")
print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

11131 exemples before dropping na

Number exemple: 10445

- training set: 8356

- test set: 2089

Number of features: p=835

Number of class: 2

class 0 : 35.1%

class 1 : 64.9%

```

In [29]: startTime = time.time()
        n_features_to_select = 20
        step = 0.05
        clf = LogisticRegression(C=1,
                                penalty='l1',
                                class_weight='balanced',
                                random_state=42)
        selector = RFE(estimator=clf, n_features_to_select=n_features_to_select, step=step)
        selector.fit(X_train, y_train)
        print(f"Optimal support of size {n_features_to_select} found in {time.time() - startTim

```

Optimal support of size 20 found in 726.4 s

```

In [30]: small_mask = selector.support_.copy()
        X_train = X_train[:,small_mask]
        X_test = X_test[:,small_mask]
        print(f"Number of features: p={X_train.shape[1]}")

```

Number of features: p=20

```

In [31]: print(f"Selected {X_train.shape[1]} features:\n")
        print(", ".join(X.columns[small_mask]))

```

Selected 20 features:

NB\_A301, NB\_D108, NB\_D703, JUSTICE, ACM8, NB\_A115, NB\_D101, NB\_F121\_NB\_COU, NB\_F116\_NB\_COU, NB\_A



#### 1.4.4 c) Using selected features to fit various models

##### Logistic regression using L1 and L2

```
In [32]: nb_value = 20 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
mean_score_l2 = np.zeros(nb_value)
C_log = np.logspace(-2.5,2,nb_value)
cv = 6 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
std_scores_l1 = np.empty(nb_value)

mean_score_l2 = np.empty(nb_value)
std_scores_l2 = np.empty(nb_value)

np.random.seed(seed=42)

startTime = time.time()

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1',
                             tol=0.01, random_state=42,
                             class_weight='balanced')
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))
    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight='balanced')
    mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))
    std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
                                                       X_train,
                                                       y_train,
                                                       cv=cv,
                                                       scoring='accuracy'))
```


```

plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:]-0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:]+0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:]-0.5*std_scores_l2[:], 'g--', label=u'+/-0.5 écart t')
plt.semilogx(C_log,mean_score_l2[:]+0.5*std_scores_l2[:],'g--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel(u"Erreur de validation croisée (%) \n(Taux moyen d'erreur de classification)")
plt.title(u"Choix de l'hyperparamètre C \npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))
plt.grid()
plt.show()
print("Détermination des paramètres optimaux en %0.1f s" % (time.time() - startTime))
print("Pénalisation l1, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l1)]))
print("Pénalisation l2, valeur optimale : C = %0.2f" % (C_log[np.argmin(mean_score_l2)]))

```



Détermination des paramètres optimaux en 28.2 s

Pénalisation l1, valeur optimale : C = 0.00

Pénalisation l2, valeur optimale : C = 0.01

```

In [33]: # Learning on full training set with optimals hyperparameters
         # and score evaluation on test set

```

```

clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l1',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Model score

- Accuracy : 70.3 %
- Precision : 73.9 % (Happy # positive class)
- Recall : 83.7 %
- F1 score : 78.5 %

In [34]: *# Learning on full training set with optimal hyperparameters  
# and score evaluation on test set*

```

clf = LogisticRegression(C=C_log[np.argmin(mean_score_l2)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)
accuracy = clf.score(X_test, y_test)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Model score

- Accuracy : 68.0 %
- Precision : 74.5 % (Happy # positive class)
- Recall : 76.8 %
- F1 score : 75.6 %

In [35]: *# Use regression coefficients to rank features*

```

clf = LogisticRegression(C=C_log[np.argmin(mean_score_l2)],
                        penalty='l2',

```

```

        random_state=42,
        class_weight='balanced')

clf.fit(X_train,y_train)
coef_l2 = abs(clf.coef_)
coef_sorted_l2 = -np.sort(-coef_l2).reshape(-1)
print(coef_sorted_l2)
features_sorded_l2 = np.argsort(-coef_l2).reshape(-1)
print(features_sorded_l2)
features_name = np.array(features)
features_name_sorted_l2 = features_name[features_sorded_l2]

clf = LogisticRegression(C=C_log[np.argmin(mean_score_l1)],
                        penalty='l2',
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train,y_train)
coef_l1 = abs(clf.coef_)
coef_sorted_l1 = -np.sort(-coef_l1).reshape(-1)
features_sorded_l1 = np.argsort(-coef_l1).reshape(-1)
features_name_sorted_l1 = features_name[features_sorded_l1]

nf = X_train.shape[1]

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
ind = np.arange(nf)    # the x locations for the groups

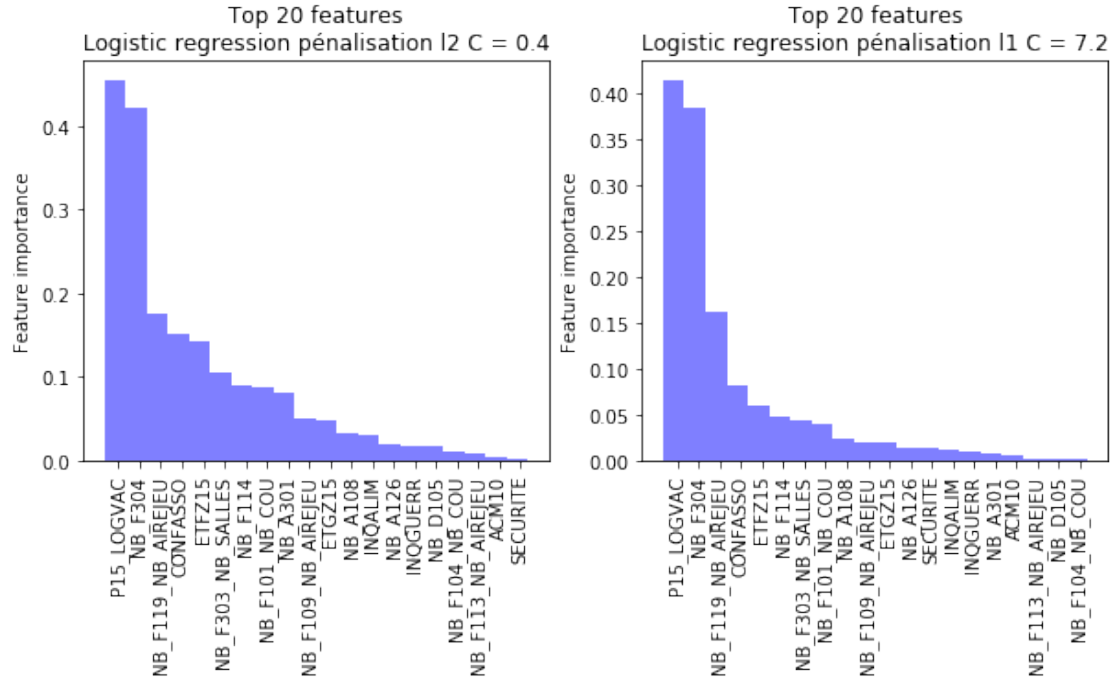
plt.subplot(1, 2, 1)
p1 = plt.bar(ind, coef_sorted_l2[0:nf], 1, color='b',alpha=0.5)
plt.ylabel('Feature importance')
plt.title(u'Top %i features\nLogistic regression pénalisation l2 C = 0.4' % nf)
plt.xticks(ind + 0.35/2.0, features_name_sorted_l2[0:nf], rotation = 90)

plt.subplot(1, 2, 2)
p1 = plt.bar(ind, coef_sorted_l1[0:nf], 1, color='b',alpha=0.5)
plt.ylabel('Feature importance')
plt.title(u'Top %i features\nLogistic regression pénalisation l1 C = 7.2' % nf)
plt.xticks(ind + 0.35/2.0, features_name_sorted_l1[0:nf], rotation = 90)

plt.show()

[ 0.4549099  0.42221413  0.17622814  0.15161577  0.14338918  0.10525766
  0.08990527  0.08769821  0.08126144  0.04908717  0.04774626  0.03318206
  0.03061024  0.02029375  0.01800904  0.01711365  0.01062747  0.00864857
  0.00407274  0.00242724]
[19 16 11  4  2  7 14  5  3 17 18  8  6  9 10  0 15 12  1 13]

```



## Random Forest

```
In [36]: startTime = time.time()
n_estimators_range = [32,64,128,256,512]
max_depth_range = [4,8,16,32,64]
param_grid = dict(n_estimators=n_estimators_range, max_depth = max_depth_range)

params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced'}
clf = RandomForestClassifier(**params)

grid = GridSearchCV(clf, scoring='accuracy', param_grid=param_grid)
grid.fit(X_train, y_train)
print(f"Determination of optimal hyperparameters in {time.time() - startTime:0.1f} s")
print(f"Optimal values are {grid.best_params_} \n\
Accuracy Score of cross validation {100*grid.best_score_:0.2f}%")

# Learning on full training set with optimals hyperparameters and score on test set
params = {'max_features' : 'sqrt', 'random_state' : 32,
          'min_samples_split' : 2, 'class_weight' : 'balanced',
          'n_estimators' : grid.best_params_['n_estimators'],
          'max_depth' : grid.best_params_['max_depth']}
clf = RandomForestClassifier(**params).fit(X_train, y_train)
clf.fit(X_train, y_train)
```

```

y_test_pred = clf.predict(X_test)

print(f"Random Forest, p={X_train.shape[1]}")
accuracy = clf.score(X_test, y_test)
f1 = f1_score(y_test, y_test_pred)
p = precision_score(y_test, y_test_pred)
r = recall_score(y_test, y_test_pred)
print(f"Model score\n- Accuracy : {accuracy*100:0.1f} %")
print(f"- Precision : {p*100:0.1f} % (Happy # positive class)")
print(f"- Recall : {r*100:0.1f} %")
print(f"- F1 score : {f1*100:0.1f} %")

```

Determination of optimal hyperparameters in 127.1 s  
 Optimal values are {'max\_depth': 4, 'n\_estimators': 256}  
 Accuracy Score of cross validation 69.17%  
 Random Forest, p=20  
 Model score  
 - Accuracy : 67.8 %  
 - Precision : 74.7 % (Happy # positive class)  
 - Recall : 76.1 %  
 - F1 score : 75.4 %

## 1.5 5) Multi class regression

### 1.5.1 a) Training and test set preparation

```

In [37]: features = df.columns.drop(['HEUREUX', "HEUREUX_CLF", "HEUREUX_REG"])
        pred = "HEUREUX"

# treating remaining missing values
df_tmp = df.loc[:,set(features) | {pred} ]
df_tmp.loc[df_tmp["HEUREUX"]==5, "HEUREUX"]=None
df_tmp = df_tmp.dropna()

X = df_tmp.loc[:,features]
y = df_tmp[pred]

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42)

scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

print(f"Number exemple: {y.shape[0]}\n- training set: \
{y_train.shape[0]}\n- test set: {y_test.shape[0]}")
print(f"Number of features: p={X_train.shape[1]}")

```

```

print(f"Number of class: {len(np.unique(y))}")
for c in np.unique(y):
    print(f"class {c:0.0f} : {100*np.sum(y==c)/len(y):0.1f}%")

```

Number exemple: 10445

- training set: 8356

- test set: 2089

Number of features: p=835

Number of class: 4

class 1 : 1.8%

class 2 : 33.3%

class 3 : 49.0%

class 4 : 15.9%

### 1.5.2 b) Feature selection

In [38]: startTime = time.time()

```

scoring='f1_macro'
step = 0.05

```

```

clf = LogisticRegression(C=1,
                        penalty='l1',
                        class_weight='balanced',
                        random_state=42)

```

```

rfecv = RFECV(estimator=clf, step=step, cv=StratifiedKFold(2),
              scoring=scoring)

```

```

rfecv.fit(X_train, y_train)

```

```

print("Optimal number of features : %d" % rfecv.n_features_)

```

```

# Plot number of features VS. cross-validation scores

```

```

plt.figure()

```

```

plt.xlabel(f"Number of recursive steps of {100*step:0.0f}% through {X_train.shape[1]} v

```

```

plt.ylabel(f"Cross validation score \n({scoring})")

```

```

plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_)

```

```

plt.show()

```

```

print(f"Détermination des features optimales en %0.1f s" % (time.time() - startTime))

```

```

//anaconda/envs/py36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: Undefined

```

```

'precision', 'predicted', average, warn_for)

```

```

//anaconda/envs/py36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: Undefined


```

```

'precision', 'predicted', average, warn_for)

```

Optimal number of features : 138



Détermination des features optimales en 8071.4 s

```
In [39]: nclf_mask = rfecv.support_.copy()
X_train = X_train[:,nclf_mask]
X_test = X_test[:,nclf_mask]
print(f"Number of features: p={X_train.shape[1]}")
```

Number of features: p=138

```
In [40]: nb_value = 20 # Nombre de valeurs testées pour l'hyperparamètre
mean_score_l1 = np.zeros(nb_value)
mean_score_l2 = np.zeros(nb_value)
C_log = np.logspace(-4,2,nb_value)
cv = 3 # V-fold, nombre de fold

mean_score_l1 = np.empty(nb_value)
std_scores_l1 = np.empty(nb_value)

mean_score_l2 = np.empty(nb_value)
std_scores_l2 = np.empty(nb_value)
```



```

np.random.seed(seed=42)

startTime = time.time()

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l1', tol=0.01, random_state=42, class_weight=
    mean_score_l1[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

    std_scores_l1[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

for i, C in enumerate(C_log):
    clf = LogisticRegression(C=C, penalty='l2', tol=0.01, random_state=42, class_weight=
    mean_score_l2[i] = 100*np.mean(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))

    std_scores_l2[i] = 100*np.std(1-cross_val_score(clf,
                                                    X_train,
                                                    y_train,
                                                    cv=cv,
                                                    scoring='f1_macro'))


plt.figure()
plt.semilogx(C_log,mean_score_l1[:],'r',linewidth=2,label='moyenne (l1)')
plt.semilogx(C_log,mean_score_l1[:] - 0.5*std_scores_l1[:],
             'r--', label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l1[:] + 0.5*std_scores_l1[:],'r--')

plt.semilogx(C_log,mean_score_l2[:],'g',linewidth=2,label='moyenne (l2)')
plt.semilogx(C_log,mean_score_l2[:] - 0.5*std_scores_l2[:], 'g--',
             label=u'+/-0.5 écart type')
plt.semilogx(C_log,mean_score_l2[:] + 0.5*std_scores_l2[:],'g--')

plt.xlabel("Valeur de pénalisation C = 1/lambda")
plt.ylabel("F1 macro de validation croisée (%)")
plt.title(u"Choix de l'hyperparamètre C\npar validation croisée \
(V-fold avec V = %s)" % (cv))
plt.legend(bbox_to_anchor=(1, 1))

```



[illegible]

Détermination des paramètres optimaux en 983.7 s  
Pénalisation l1, valeur optimale : C = 0.0001  
Pénalisation l2, valeur optimale : C = 0.0001

```
In [41]: # Learning on full training set with optimals hyperparameters and score on test set
clf = LogisticRegression(C=C_log[np.argmax(mean_score_l2)],
                        penalty='l2',
                        tol=0.01,
                        random_state=42,
                        class_weight='balanced')

clf.fit(X_train, y_train)
y_test_pred = clf.predict(X_test)

In [42]: def plot_confusion_matrix(cm, classes,
                                normalize=False,
                                title='Confusion matrix',
                                cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```
In [43]: # Model evaluation
```

```

class_names = ["Jamais",
               "Occasionnellement",
               "Assez souvent",
               "Très souvent" ]

f1_scores = f1_score(y_test, y_test_pred, labels = [1,2,3,4], average=None)
for i,c in enumerate(class_names):
    print(f"f1 score class '{c}' : {100*f1_scores[i]:0.1f}%")
accuracy = clf.score(X_test, y_test)
f1_macro = f1_score(y_test, y_test_pred, average='macro')
f1_weighted = f1_score(y_test, y_test_pred, average='weighted')
print(f"Average scores :\nf1 macro : {f1_macro*100:0.4f} %\n\
f1 weighted : {f1_weighted*100:0.4f} %\naccuracy : {accuracy*100:0.4f} %")

# Compute confusion matrix
cnf_matrix = confusion_matrix(y_test, y_test_pred)
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()

f1 score class 'Jamais' : 11.7%
f1 score class 'Occasionnellement' : 55.5%
f1 score class 'Assez souvent' : 53.9%
f1 score class 'Très souvent' : 37.1%
Average scores :
f1 macro : 39.5539 %
f1 weighted : 50.8346 %
accuracy : 50.3112 %
Confusion matrix, without normalization
[[ 7 28  1  6]
 [33 414 179 69]
 [30 302 493 182]
 [ 8 52 148 137]]
Normalized confusion matrix
[[ 0.17  0.67  0.02  0.14]
 [ 0.05  0.6  0.26  0.1 ]
 [ 0.03  0.3  0.49  0.18]
 [ 0.02  0.15  0.43  0.4 ]]

```

